NISTIR XXXX Draft

Ongoing Face Recognition Vendor Test (FRVT)

Part 1: Verification

Patrick Grother
Mei Ngan
Kayee Hanaoka
Information Access Division
Information Technology Laboratory

This publication is available free of charge from:
https://www.nist.gov/programs-projects/face-recognition-vendor-test-frvt-ongoing

National Institute of Standards and Technology
U.S. Department of Commerce

DISCLAIMER

Specific hardware and software products identified in this report were used in order to perform the evaluations described in this document. In no case does identification of any commercial product, trade name, or vendor, imply recommendation or endorsement by the National Institute of Standards and Technology, nor does it imply that the products and equipment identified are necessarily the best available for the purpose.

ABOUT THIS REPORT

This report is a draft NIST Interagency Report, and is open for comment. It documents the verification-track of the ongoing Face Recognition Vendor Test. The report will be updated continuously as new algorithms are evaluated, as new datasets are added, and as new analyses are included. Comments and suggestions should be directed to frvt@nist.gov.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Contents

DISCLAIMER 1

1 NEWS 5

2 CHANGELOG 6

3 METRICS 11
3.1 CORE ACCURACY . . . 11

4 DATASETS 12
4.1 CHILD EXPLOITATION IMAGES . . . 12
4.2 VISA IMAGES . . . 12
4.3 MUGSHOT IMAGES . . . 12
4.4 SELFIE IMAGES . . . 13
4.5 WEBCAM IMAGES . . . 13
4.6 WILD IMAGES . . . 14

5 RESULTS 14
5.1 TEST GOALS . . . 14
5.2 TEST DESIGN . . . 14
5.3 FAILURE TO ENROL . . . 17
5.4 RECOGNITION ACCURACY . . . 18
5.5 GENUINE DISTRIBUTION STABILITY . . . 33
5.5.1 EFFECT OF BIRTH PLACE ON THE GENUINE DISTRIBUTION . . . 33
5.5.2 EFFECT OF AGE ON GENUINE SUBJECTS . . . 35
5.6 IMPOSTOR DISTRIBUTION STABILITY . . . 37
5.6.1 EFFECT OF BIRTH PLACE ON THE IMPOSTOR DISTRIBUTION . . . 37
5.6.2 EFFECT OF AGE ON IMPOSTORS . . . 107

List of Tables

1 ALGORITHM SUMMARY . . . 7
2 FALSE NON-MATCH RATE . . . 8
3 FAILURE TO ENROL RATES . . . 17

List of Figures

1 PERFORMANCE SUMMARY: FNMR VS. TEMPLATE SIZE TRADEOFF . . . 9
2 PERFORMANCE SUMMARY: FNMR VS. TEMPLATE TIME TRADEOFF . . . 10
3 EXAMPLE IMAGES . . . 14
(A) VISA . . . 14
(B) MUGSHOT . . . 14
(C) WEBCAM . . . 14
(D) SELFIE . . . 14
(E) WILD . . . 14
4 ERROR TRADEOFF CHARACTERISTIC: VISA IMAGES . . . 19
5 ERROR TRADEOFF CHARACTERISTIC: VISA IMAGES . . . 20
6 ERROR TRADEOFF CHARACTERISTIC: MUGSHOT IMAGES . . . 21
7 ERROR TRADEOFF CHARACTERISTIC: SELFIE IMAGES . . . 22
8 ERROR TRADEOFF CHARACTERISTIC: SELFIE IMAGES . . . 23

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

9 ERROR TRADEOFF CHARACTERISTIC: WILD IMAGES . . . 24
10 ERROR TRADEOFF CHARACTERISTICS: CHILD EXPLOITATION IMAGES . . . 25
11 ERROR TRADEOFF CHARACTERISTICS: CHILD EXPLOITATION IMAGES . . . 26
12 SEX AND RACE EFFECTS: MUGSHOT IMAGES . . . 27
13 SEX EFFECTS: VISA IMAGES . . . 28
14 ERROR TRADEOFF CHARACTERISTIC: WILD IMAGES . . . 29
15 FALSE MATCH RATE CALIBRATION: VISA IMAGES . . . 30
16 FALSE MATCH RATE CONCENTRATION: VISA IMAGES . . . 31
17 FALSE MATCH RATE CALIBRATION: MUGSHOT IMAGES . . . 32
18 EFFECT OF COUNTRY OF BIRTH ON FNMR . . . 34
19 EFFECT OF SUBJECT AGE ON FNMR . . . 36
20 IMPOSTOR DISTRIBUTION SHIFTS FOR SELECT COUNTRY PAIRS . . . 39
21 ALGORITHM 3DIVI-000 CROSS REGION FMR . . . 40
22 ALGORITHM 3DIVI-001 CROSS REGION FMR . . . 41
23 ALGORITHM CAMVI-001 CROSS REGION FMR . . . 42
24 ALGORITHM CYBEREXTRUDER-001 CROSS REGION FMR . . . 43
25 ALGORITHM DERMALOG-002 CROSS REGION FMR . . . 44
26 ALGORITHM DERMALOG-003 CROSS REGION FMR . . . 45
27 ALGORITHM DIGITALBARRIERS-000 CROSS REGION FMR . . . 46
28 ALGORITHM DIGITALBARRIERS-001 CROSS REGION FMR . . . 47
29 ALGORITHM ID3-001 CROSS REGION FMR . . . 48
30 ALGORITHM ID3-002 CROSS REGION FMR . . . 49
31 ALGORITHM ISITYOU-000 CROSS REGION FMR . . . 50
32 ALGORITHM MORPHO-000 CROSS REGION FMR . . . 51
33 ALGORITHM MORPHO-001 CROSS REGION FMR . . . 52
34 ALGORITHM NEUROTECHNOLOGY-000 CROSS REGION FMR . . . 53
35 ALGORITHM NEUROTECHNOLOGY-001 CROSS REGION FMR . . . 54
36 ALGORITHM NOBLIS-000 CROSS REGION FMR . . . 55
37 ALGORITHM NTECHLAB-001 CROSS REGION FMR . . . 56
38 ALGORITHM NTECHLAB-002 CROSS REGION FMR . . . 57
39 ALGORITHM RANKONE-000 CROSS REGION FMR . . . 58
40 ALGORITHM RANKONE-002 CROSS REGION FMR . . . 59
41 ALGORITHM SAMTECH-000 CROSS REGION FMR . . . 60
42 ALGORITHM TONGYITRANS-001 CROSS REGION FMR . . . 61
43 ALGORITHM TONGYITRANS-002 CROSS REGION FMR . . . 62
44 ALGORITHM VCOG-001 CROSS REGION FMR . . . 63
45 ALGORITHM VCOG-002 CROSS REGION FMR . . . 64
46 ALGORITHM VIGILANTSOLUTIONS-000 CROSS REGION FMR . . . 65
47 ALGORITHM VIGILANTSOLUTIONS-001 CROSS REGION FMR . . . 66
48 ALGORITHM VISIONLABS-001 CROSS REGION FMR . . . 67
49 ALGORITHM VISIONLABS-002 CROSS REGION FMR . . . 68
50 ALGORITHM VOCORD-002 CROSS REGION FMR . . . 69
51 ALGORITHM YISHENG-000 CROSS REGION FMR . . . 70
52 ALGORITHM YISHENG-001 CROSS REGION FMR . . . 71
53 ALGORITHM YITU-000 CROSS REGION FMR . . . 72
54 ALGORITHM 3DIVI-000 CROSS COUNTRY FMR . . . 73
55 ALGORITHM 3DIVI-001 CROSS COUNTRY FMR . . . 74
56 ALGORITHM CAMVI-001 CROSS COUNTRY FMR . . . 75
57 ALGORITHM CYBEREXTRUDER-001 CROSS COUNTRY FMR . . . 76
58 ALGORITHM DERMALOG-002 CROSS COUNTRY FMR . . . 77
59 ALGORITHM DERMALOG-003 CROSS COUNTRY FMR . . . 78
60 ALGORITHM DIGITALBARRIERS-000 CROSS COUNTRY FMR . . . 79
61 ALGORITHM DIGITALBARRIERS-001 CROSS COUNTRY FMR . . . 80
62 ALGORITHM ID3-001 CROSS COUNTRY FMR . . . 81
63 ALGORITHM ID3-002 CROSS COUNTRY FMR . . . 82
64 ALGORITHM ISITYOU-000 CROSS COUNTRY FMR . . . 83
65 ALGORITHM MORPHO-000 CROSS COUNTRY FMR . . . 84

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

66 ALGORITHM MORPHO-001 CROSS COUNTRY FMR . . . 85
67 ALGORITHM NEUROTECHNOLOGY-000 CROSS COUNTRY FMR . . . 86
68 ALGORITHM NEUROTECHNOLOGY-001 CROSS COUNTRY FMR . . . 87
69 ALGORITHM NOBLIS-000 CROSS COUNTRY FMR . . . 88
70 ALGORITHM NTECHLAB-001 CROSS COUNTRY FMR . . . 89
71 ALGORITHM NTECHLAB-002 CROSS COUNTRY FMR . . . 90
72 ALGORITHM RANKONE-000 CROSS COUNTRY FMR . . . 91
73 ALGORITHM RANKONE-002 CROSS COUNTRY FMR . . . 92
74 ALGORITHM SAMTECH-000 CROSS COUNTRY FMR . . . 93
75 ALGORITHM TONGYITRANS-001 CROSS COUNTRY FMR . . . 94
76 ALGORITHM TONGYITRANS-002 CROSS COUNTRY FMR . . . 95
77 ALGORITHM VCOG-001 CROSS COUNTRY FMR . . . 96
78 ALGORITHM VCOG-002 CROSS COUNTRY FMR . . . 97
79 ALGORITHM VIGILANTSOLUTIONS-000 CROSS COUNTRY FMR . . . 98
80 ALGORITHM VIGILANTSOLUTIONS-001 CROSS COUNTRY FMR . . . 99
81 ALGORITHM VISIONLABS-001 CROSS COUNTRY FMR . . . 100
82 ALGORITHM VISIONLABS-002 CROSS COUNTRY FMR . . . 101
83 ALGORITHM VOCORD-002 CROSS COUNTRY FMR . . . 102
84 ALGORITHM YISHENG-000 CROSS COUNTRY FMR . . . 103
85 ALGORITHM YISHENG-001 CROSS COUNTRY FMR . . . 104
86 ALGORITHM YITU-000 CROSS COUNTRY FMR . . . 105
87 IMPOSTOR COUNTS FOR CROSS COUNTRY FMR CALCULATIONS . . . 106
88 ALGORITHM 3DIVI-000 CROSS AGE FMR . . . 108
89 ALGORITHM 3DIVI-001 CROSS AGE FMR . . . 109
90 ALGORITHM CAMVI-001 CROSS AGE FMR . . . 110
91 ALGORITHM CYBEREXTRUDER-001 CROSS AGE FMR . . . 111
92 ALGORITHM DERMALOG-002 CROSS AGE FMR . . . 112
93 ALGORITHM DERMALOG-003 CROSS AGE FMR . . . 113
94 ALGORITHM DIGITALBARRIERS-000 CROSS AGE FMR . . . 114
95 ALGORITHM DIGITALBARRIERS-001 CROSS AGE FMR . . . 115
96 ALGORITHM ID3-001 CROSS AGE FMR . . . 116
97 ALGORITHM ID3-002 CROSS AGE FMR . . . 117
98 ALGORITHM ISITYOU-000 CROSS AGE FMR . . . 118
99 ALGORITHM MORPHO-000 CROSS AGE FMR . . . 119
100 ALGORITHM MORPHO-001 CROSS AGE FMR . . . 120
101 ALGORITHM NEUROTECHNOLOGY-000 CROSS AGE FMR . . . 121
102 ALGORITHM NEUROTECHNOLOGY-001 CROSS AGE FMR . . . 122
103 ALGORITHM NOBLIS-000 CROSS AGE FMR . . . 123
104 ALGORITHM NTECHLAB-001 CROSS AGE FMR . . . 124
105 ALGORITHM NTECHLAB-002 CROSS AGE FMR . . . 125
106 ALGORITHM RANKONE-000 CROSS AGE FMR . . . 126
107 ALGORITHM RANKONE-002 CROSS AGE FMR . . . 127
108 ALGORITHM SAMTECH-000 CROSS AGE FMR . . . 128
109 ALGORITHM TONGYITRANS-001 CROSS AGE FMR . . . 129
110 ALGORITHM TONGYITRANS-002 CROSS AGE FMR . . . 130
111 ALGORITHM VCOG-001 CROSS AGE FMR . . . 131
112 ALGORITHM VCOG-002 CROSS AGE FMR . . . 132
113 ALGORITHM VIGILANTSOLUTIONS-000 CROSS AGE FMR . . . 133
114 ALGORITHM VIGILANTSOLUTIONS-001 CROSS AGE FMR . . . 134
115 ALGORITHM VISIONLABS-001 CROSS AGE FMR . . . 135
116 ALGORITHM VISIONLABS-002 CROSS AGE FMR . . . 136
117 ALGORITHM VOCORD-002 CROSS AGE FMR . . . 137
118 ALGORITHM YISHENG-000 CROSS AGE FMR . . . 138
119 ALGORITHM YISHENG-001 CROSS AGE FMR . . . 139
120 ALGORITHM YITU-000 CROSS AGE FMR . . . 140

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

1 News

2017-10-03

- In 2018, NIST will run an evaluation of 1:N identification algorithms.
- The test will use N well in excess of 10^7.
- The test will proceed with a fixed participation deadline, likely in February 2018.
- A draft of the API and the concept of operations will be released in October 2017.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

2 Changelog

2017-10-03

- ▷ Added results for algorithms from 3DiVi, Camvi, Idemia, Noblis, N-TechLab, and Visionlabs.
- ▷ Added partial results for two algorithms from Zhuhai Yisheng.
- ▷ The ntechlab-000 algorithm has been retired - FRVT lists only two algorithms per organization.
- ▷ Corrected fixed FMR operating point in the legends of some DET plots.

2017-08-22

- ▷ Added results for three additional algorithms, rankone-002, neurotechnology-001, and itmo-002.
- ▷ The algorithms dermalog-001 and rank-001 have been retired - FRVT lists only two algorithms per organization
- ▷ The algorithm tupel-001 has been retired as it is not operable on all datasets
- ▷ Clarified the tradeoff Figures 1 and 2 plot only genuine comparison durations.
- ▷ Corrected image type label in section 4.5

2017-08-07

- ▷ Added results for 5 new algorithms
- ▷ Added Figure 3 giving simulated example images.
- ▷ Added Figure 1 showing an alternative view of the same tradeoff data in Figure 2
- ▷ Added Figure 5 showing accuracy on visa images just for low FMR.
- ▷ Added Figure 20 showing impostor distribution shifts from certain country pairs. Section 5.6.1 in this and prior reports documents high false match rates for individuals from certain countries. That effect, however, is often not confined to anomalously high impostor scores in the tails of the distribution, but arises from systematic shifts of the whole distribution. These shifts sometimes reach 2σ.

2017-07-29

- ▷ Added results for 8 new algorithms
- ▷ Added results for a child-exploitation dataset
- ▷ Added Table 2 a standalone tabulation of false non-match rates
- ▷ We have received additional CPU algorithms - Results should appear August 4, 2017
- ▷ We have received additional GPU algorithms - Results to appear as computational resources are released from the Face Recognition Prize Challenge

2017-06-19

- ▷ Added five new algorithms, three of which remain in-process
- ▷ Added results for a "wild" dataset of images similar to non-cooperative photojournalism images
- ▷ Added Table 3 a standalone tabulation of failure to enrol rates
- ▷ Added Fig. 2 showing tradeoff between FNMR, template size, template generation time, and match duration.
- ▷ Added Fig. 16 showing how FMR is concentrated in certain images.
- ▷ Restated cross-region false match rates at nominal FMR = 0.0001 instead of 0.001
- ▷ Improved DET legends.

	Developer	Short	Seq.	Validation	Config[1]	Template		GPU	Comparison Time (ns)[3]	
	Name	Name	Num.	Date	Data (KB)	Size (B)	Time (ms)[2]		Genuine	Impostor
1	3DiVi	3divi	000	2017-03-16	169360	$^{10}512 \pm 0$	$^{19}285 \pm 52$	Yes	$^{1}378 \pm 20$	$^{2}375 \pm 19$
2	3DiVi	3divi	001	2017-06-22	190867	$^{28}4096 \pm 0$	$^{18}274 \pm 47$	Yes	$^{4}636 \pm 19$	$^{5}634 \pm 16$
3	Ayonix	ayonix	000	2017-06-22	58505	$^{14}1036 \pm 0$	$^{1}18 \pm 2$	No	$^{3}621 \pm 23$	$^{4}620 \pm 26$
4	Camvi Technologies	camvitech	001	2017-09-13	118759	$^{13}1024 \pm 0$	$^{14}181 \pm 6$	No	$^{2}481 \pm 16$	$^{3}487 \pm 23$
5	Cyberextruder	cyberex	001	2017-08-02	121211	$^{7}256 \pm 0$	$^{35}893 \pm 25$	No	$^{12}1083 \pm 16$	$^{14}1079 \pm 19$
6	Dermalog	dermalog	002	2017-02-22	0	$^{16}1043 \pm 0$	$^{4}81 \pm 0$	No	$^{25}22169 \pm 114$	$^{28}22105 \pm 146$
7	Dermalog	dermalog	003	2017-07-10	0	$^{15}1043 \pm 0$	$^{12}121 \pm 22$	No	$^{26}22957 \pm 93$	$^{29}22808 \pm 131$
8	Digital Barriers	barriers	000	2017-05-31	157794	$^{22}2056 \pm 0$	$^{8}104 \pm 0$	No	$^{21}13232 \pm 166$	$^{24}13226 \pm 146$
9	Digital Barriers	barriers	001	2017-07-20	236934	$^{24}2056 \pm 0$	$^{20}294 \pm 1$	No	$^{20}12311 \pm 164$	$^{22}12347 \pm 197$
10	ID3 Technology	id3	001	2017-08-04	225574	$^{12}520 \pm 0$	$^{15}238 \pm 19$	No	$^{11}1058 \pm 32$	$^{13}1049 \pm 28$
11	ID3 Technology	id3	002	2017-08-04	225574	$^{11}520 \pm 0$	$^{25}482 \pm 34$	No	$^{13}1100 \pm 59$	$^{12}1048 \pm 32$
12	Innovatrics	innova	000	2017-07-25	0	$^{4}146 \pm 0$	$^{26}578 \pm 5$	No	$^{17}4964 \pm 63$	$^{18}4665 \pm 262$
13	Innovatrics	innova	001	2017-07-25	0	$^{9}284 \pm 0$	$^{30}645 \pm 5$	No	$^{18}5506 \pm 131$	$^{19}4975 \pm 308$
14	Is It You	isityou	000	2017-06-26	48010	$^{36}19200 \pm 0$	$^{11}113 \pm 5$	No	$^{37}237517 \pm 1318$	$^{37}237374 \pm 1279$
15	ITMO University	itmo	001	2017-06-12	1923215	$^{38}37997 \pm 0$	$^{39}959 \pm 18$	No	$^{27}28901 \pm 1492$	$^{30}27517 \pm 186$
16	ITMO University	itmo	002	2017-08-07	1923215	$^{31}4162 \pm 0$	$^{28}611 \pm 17$	No	$^{19}7423 \pm 96$	$^{20}7451 \pm 94$
17	Idemia	morpho	000	2017-07-11	100806	$^{1}116 \pm 0$	$^{10}109 \pm 1$	No	$^{9}993 \pm 31$	$^{10}1000 \pm 34$
18	Idemia	morpho	001	2017-07-10	100806	$^{2}116 \pm 0$	$^{2}41 \pm 1$	No	$^{10}1030 \pm 38$	$^{11}1042 \pm 37$
19	Neurotechnology	neurotech	000	2017-03-22	62129	$^{35}7148 \pm 0$	$^{29}611 \pm 48$	No	$^{34}74288 \pm 2194$	$^{34}72879 \pm 2640$
20	Neurotechnology	neurotech	001	2017-08-07	280771	$^{27}2718 \pm 0$	$^{33}881 \pm 46$	No	$^{33}69356 \pm 684$	$^{33}69140 \pm 579$
21	Noblis	noblis	000	2017-07-05	713573	$^{17}1061 \pm 0$	$^{13}174 \pm 20$	Yes	$^{14}1343 \pm 45$	$^{16}1373 \pm 56$
22	N-Tech Lab	ntech	001	2017-05-10	691296	$^{34}6744 \pm 1$	$^{27}587 \pm 11$	No	$^{31}67692 \pm 833$	$^{32}67486 \pm 244$
23	N-Tech Lab	ntech	002	2017-08-23	894169	$^{33}6744 \pm 1$	$^{22}330 \pm 11$	No	$^{35}82508 \pm 213$	$^{35}82524 \pm 220$
24	Rank One Computing	rankone	000	2017-03-21	0	$^{3}144 \pm 0$	$^{5}82 \pm 9$	No	$^{30}39932 \pm 468$	$^{21}8722 \pm 171$
25	Rank One Computing	rankone	002	2017-08-18	0	$^{6}224 \pm 0$	$^{3}75 \pm 1$	No	$^{32}69113 \pm 3802$	$^{1}369 \pm 26$
26	Samtech InfoNet Limited	samtech	000	2017-05-02	109774	$^{23}2056 \pm 0$	$^{16}262 \pm 2$	No	$^{16}4550 \pm 26$	$^{17}4541 \pm 28$
27	TongYi Transportation Technology	tongyi	001	2017-04-01	625339	$^{25}2058 \pm 0$	$^{21}310 \pm 20$	No	$^{23}17769 \pm 74$	$^{26}17750 \pm 63$
28	TongYi Transportation Technology	tongyi	002	2017-07-15	625336	$^{26}2058 \pm 0$	$^{23}356 \pm 35$	No	$^{28}29816 \pm 281$	$^{27}17799 \pm 127$
29	VCognition	vcog	001	2017-03-28	86103	$^{29}4126 \pm 0$	$^{9}108 \pm 17$	Yes	$^{22}16320 \pm 197$	$^{25}16426 \pm 425$
30	VCognition	vcog	002	2017-06-12	3229434	$^{39}61504 \pm 5$	$^{24}357 \pm 25$	No	$^{38}296154 \pm 3077$	$^{38}296436 \pm 4183$
31	Vigilant Solutions	vigilant	000	2017-03-30	352218	$^{37}31540 \pm 0$	$^{34}884 \pm 23$	No	$^{24}18201 \pm 94$	$^{23}13030 \pm 83$
32	Vigilant Solutions	vigilant	001	2017-06-13	344685	$^{19}1544 \pm 0$	$^{37}921 \pm 2$	No	$^{5}644 \pm 13$	$^{6}649 \pm 16$
33	VisionLabs	visionlabs	001	2017-06-12	343661	$^{5}204 \pm 0$	$^{38}943 \pm 8$	No	$^{15}1395 \pm 45$	$^{15}1148 \pm 53$
34	VisionLabs	visionlabs	002	2017-09-08	591936	$^{8}264 \pm 0$	$^{17}265 \pm 8$	No	$^{7}796 \pm 35$	$^{8}799 \pm 31$
35	Vocord	vocord	001	2017-04-21	616989	$^{32}6194 \pm 0$	$^{36}908 \pm 16$	No	$^{39}1094730 \pm 64282$	$^{39}1107193 \pm 66523$
36	Vocord	vocord	002	2017-06-07	918292	$^{18}1330 \pm 0$	$^{32}782 \pm 36$	Yes	$^{36}83063 \pm 517$	$^{36}83072 \pm 714$
37	Zhuhai Yisheng Electronics Technology	yisheng	000	2017-08-17	122704	$^{21}2048 \pm 0$	$^{7}103 \pm 1$	Yes	$^{6}790 \pm 23$	$^{7}789 \pm 20$
38	Zhuhai Yisheng Electronics Technology	yisheng	001	2017-08-17	120112	$^{20}2048 \pm 0$	$^{6}103 \pm 1$	Yes	$^{8}906 \pm 31$	$^{9}905 \pm 25$
39	Shanghai Yitu Technology	yitu	000	2017-05-23	2211068	$^{30}4130 \pm 0$	$^{31}672 \pm 2$	No	$^{29}35352 \pm 114$	$^{31}37848 \pm 1773$

Notes	
1	The size of configuration data does not capture static data included in the libraries. We do not include the size of the libraries because some algorithms include common ancilliary libraries for image processing (e.g. openCV) or numerical computation (e.g. blas).
2	The median template creation times are measured on Intel®Xeon®CPU E5-2630 v4 @ 2.20GHz processors or, in the case of GPU-enabled implementations, NVidia Tesla K40.
3	The median comparison durations, in nanoseconds, are estimated using std::chrono::high_resolution_clock which on the machine in (2) counts clock ticks of duration 1ns. Precision is somewhat worse than that however. The $\pm$ value is the median absolute deviation times 1.48 for Normal consistency.

Table 1: Summary of algorithms and properties included in this report. The red superscripts give ranking for the quantity in that column.

		FALSE NON-MATCH RATE (FNMR)													
	Algorithm	CONSTRAINED, COOPERATIVE								LESS CONSTRAINED, NON-COOPERATIVE					
	Name	VISA		VISA		MUGSHOT		WEBCAM		SELFIE		WILD		CHILD EXP	
	FMR	1E-06		0.0001		0.0001		0.0001		0.0001		0.0001		0.01	
1	3divi-000	0.133	8	0.029	8	0.037	13	0.002	8	0.055	14	0.547	6	0.553	5
2	3divi-001	0.154	11	0.020	3	0.030	9	-	36	-	36	0.492	5	0.549	4
3	ayonix-000	0.487	32	0.230	35	0.309	31	0.172	29	0.360	29	0.807	24	0.843	31
4	camvi-001	0.538	34	0.183	33	0.323	32	0.106	27	0.268	28	0.743	21	0.763	22
5	cyberextruder-001	0.255	22	0.076	25	0.165	25	0.029	20	0.144	24	0.853	27	0.780	26
6	dermalog-002	0.315	27	0.122	28	0.241	29	0.109	28	0.179	26	0.709	15	0.985	38
7	dermalog-003	0.280	25	0.112	27	0.202	28	0.041	24	0.115	23	0.693	12	0.845	33
8	digitalbarriers-000	0.463	31	0.161	30	0.184	27	0.045	25	0.170	25	0.741	20	0.771	24
9	digitalbarriers-001	0.502	33	0.155	29	0.041	16	0.029	21	0.115	22	0.678	10	0.731	14
10	id3-001	0.250	21	0.063	23	0.036	12	0.002	9	0.040	10	0.765	22	0.688	11
11	id3-002	0.239	20	0.057	22	0.032	11	0.003	10	0.037	8	0.810	25	0.686	9
12	innovatrics-000	0.191	15	0.034	14	0.046	20	0.001	5	0.040	9	0.720	16	0.852	35
13	innovatrics-001	0.183	14	0.034	13	0.043	17	0.001	6	0.043	11	0.643	8	0.815	30
14	isityou-000	0.703	37	0.414	37	0.680	35	0.690	33	-	34	1.000	37	-	39
15	itmo-001	0.441	29	0.174	32	0.256	30	0.069	26	0.199	27	0.932	34	0.797	29
16	itmo-002	0.287	26	0.050	20	0.047	21	0.036	22	0.069	20	0.892	30	0.783	27
17	morpho-000	0.134	9	0.026	7	0.028	7	0.007	13	0.012	2	0.893	31	0.846	34
18	morpho-001	0.277	24	0.033	12	0.028	6	-	39	-	39	0.931	33	0.901	37
19	neurotechnology-000	0.237	19	0.052	21	0.062	22	0.005	12	0.052	12	0.943	35	0.845	32
20	neurotechnology-001	0.222	18	0.044	18	0.046	19	0.001	7	0.017	6	0.817	26	0.533	3
21	noblis-000	0.542	35	0.212	34	0.349	33	-	37	-	37	0.875	28	0.757	20
22	ntechlab-001	0.083	7	0.025	6	0.030	8	0.003	11	0.014	4	0.319	1	0.472	2
23	ntechlab-002	0.065	4	0.021	4	0.023	4	-	34	-	33	0.324	2	0.459	1
24	rankone-000	0.276	23	0.071	24	0.177	26	0.021	18	0.092	21	0.723	17	0.787	28
25	rankone-002	0.217	17	0.049	19	0.032	10	0.000	2	0.052	13	0.705	14	0.778	25
26	samtech-000	0.443	30	0.161	31	0.044	18	0.021	19	0.063	18	0.878	29	0.765	23
27	tongyitrans-001	0.072	6	0.038	17	0.041	15	0.009	14	0.063	16	0.704	13	0.743	15
28	tongyitrans-002	0.066	5	0.030	9	0.039	14	0.010	15	0.063	17	0.725	18	0.746	16
29	vcog-001	0.892	38	0.409	36	-	38	0.302	30	0.427	30	-	38	0.686	10
30	vcog-002	0.903	39	0.504	39	0.692	36	0.559	32	0.666	32	0.778	23	0.752	18
31	vigilantsolutions-000	0.688	36	0.415	38	0.595	34	0.401	31	0.643	31	0.915	32	0.894	36
32	vigilantsolutions-001	0.348	28	0.105	26	0.101	24	0.016	17	0.061	15	0.729	19	0.730	13
33	visionlabs-001	0.180	13	0.030	10	0.024	5	0.001	3	0.014	5	0.591	7	0.561	7
34	visionlabs-002	0.051	3	0.014	2	0.020	3	0.001	4	0.017	7	0.397	3	0.556	6
35	vocord-001	0.141	10	0.035	15	0.063	23	0.036	23	0.069	19	0.654	9	0.695	12
36	vocord-002	0.034	2	0.013	1	0.019	2	0.011	16	0.012	3	0.948	36	0.762	21
37	yisheng-000	0.199	16	0.037	16	-	37	-	35	-	35	0.686	11	0.749	17
38	yisheng-001	0.160	12	0.032	11	-	39	-	38	-	38	-	39	0.756	19
39	yitu-000	0.033	1	0.021	5	0.017	1	0.000	1	0.012	1	0.431	4	0.586	8

Table 2: FNMR is the proportion of mated comparisons below a threshold set to achieve the FMR given in the header on the fourth row. FMR is the proportion of impostor comparisons at or above that threshold. Note that the webcam and selfie values apply to images collected on the same day, and that will often yield optimisticly low FNMR values.

Figure 1: The points show false non-match rates (FNMR) versus the duration of the template generation operation. FNMR is the geometric mean of FNMR values for visa and mugshot images (from Figs. 4 and 6) at a false match rate (FMR) of 0.0001. Template generation time is a median estimated over 640 x 480 pixel portraits. The size of the points encodes template generation time - which spans one order of magnitude. The color of the points encodes one-to-one genuine template comparison duration - which span three orders of magnitude.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 2: The points show false non-match rates (FNMR) versus the duration of the template generation operation. FNMR is the geometric mean of FNMR values for visa and mugshot images (from Figs. 4 and 6) at a false match rate (FMR) of 0.0001. Template generation time is a median estimated over 640 x 480 pixel portraits. The size of the points encodes template size - which span two orders of magnitude. The color of the points encodes one-to-one genuine template comparison duration - which span three orders of magnitude.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

3 Metrics

3.1 Core accuracy

Given a vector of N genuine scores, u, the false non-match rate (FNMR) is computed as the proportion below some threshold, T:

$$\text{FNMR}(T) = 1 - \frac{1}{N}\sum_{i=1}^{N} H(u_i - T) \tag{1}$$

where $H(x)$ is the unit step function, and $H(0)$ taken to be 1.

Similarly, given a vector of N impostor scores, v, the false match rate (FMR) is computed as the proportion above T:

$$\text{FMR}(T) = \frac{1}{N}\sum_{i=1}^{N} H(v_i - T) \tag{2}$$

The threshold, T, can take on any value. We typically generate a set of thresholds from quantiles of the observed impostor scores, v, as follows. Given some interesting false match rate range, $[\text{FMR}_L, \text{FMR}_U]$, we form a vector of K thresholds corresponding to FMR measurements evenly spaced on a logarithmic scale

$$T_k = Q_v(1 - \text{FMR}_k) \tag{3}$$

where Q is the quantile function, and FMR_k comes from

$$\log_{10}\text{FMR}_k = \log_{10}\text{FMR}_L + \frac{k}{K}\,[\log_{10}\text{FMR}_U - \log_{10}\text{FMR}_L] \tag{4}$$

Error tradeoff characteristics are plots of FNMR(T) vs. FMR(T). These are plotted with $\text{FMR}_U \to 1$ and FMR_L as low as is sustained by the number of impostor comparisons, N. This is somewhat higher than the "rule of three" limit $3/N$ because samples are not independent, due to re-use of images.

4 Datasets

4.1 Child exploitation images

- ▷ The number of images is O(10^4).
- ▷ The number of subjects is O(10^3).
- ▷ The number of subjects with two images O(10^3).
- ▷ The images are operational. They are taken from ongoing investigations of child exploitation crimes. The images are arbitrarily unconstrained. Pose varies considerably around all three axes, including subject lying down. Resolution varies very widely. Faces can be occluded by other objects, including hair and hands. Lighting varies, although the images are intended for human viewing. Mis-focus is rare. Images are given to the algorithm without any cropping; faces may occupy widely varying areas.
- ▷ The images are usually large from contemporary cameras. The mean interocular distance (IOD) is 70 pixels.
- ▷ The images are of subjects from several countries, due to the global production of this imagery.
- ▷ The images are of children, from infancy to late adolescence.
- ▷ All of the images are live capture, none are scanned. Many have been cropped.
- ▷ When these images are input to the algorithm, they are labelled as being of type "EXPLOITATION" - see Table 4 of the FRVT API.

4.2 Visa images

- ▷ The number of images is O(10^5).
- ▷ The number of subjects is O(10^5).
- ▷ The number of subjects with two images O(10^4).
- ▷ The images have geometry in reasonable conformance with the ISO/IEC 19794-5 Full Frontal image type. Pose is generally excellent.
- ▷ The images are of size 252x300 pixels. The mean interocular distance (IOD) is 69 pixels.
- ▷ The images are of subjects from greater than 100 countries, with significant imbalance due to visa issuance patterns.
- ▷ The images are of subjects of all ages, including children, again with imbalance due to visa issuance demand.
- ▷ Many of the images are live capture. A substantial number of the images are photographs of paper photographs.
- ▷ When these images are input to the algorithm, they are labelled as being of type "ISO" - see Table 4 of the FRVT API.

4.3 Mugshot images

- ▷ The number of images is O(10^6).
- ▷ The number of subjects is O(10^5).
- ▷ The number of subjects with two images O(10^5).

- ▷ The images have geometry in reasonable conformance with the ISO/IEC 19794-5 Full Frontal image type.
- ▷ The images are of variable sizes. The median IOD is 104 pixels. The mean IOD is 123 pixels.
- ▷ The images are of subjects from the United States.
- ▷ The images are of adults.
- ▷ The images are all live capture.
- ▷ When these images are input to the algorithm, they are labelled as being of type ”mugshot” - see Table 4 of the FRVT API.

4.4 Selfie images

- ▷ The number of images is below 500.
- ▷ The number of subjects is below 500.
- ▷ All subjects have a selfie image, and a portrait image.
- ▷ The portrait images are in reasonable conformance with the ISO/IEC 19794-5 Full Frontal image type.
- ▷ The selfie images vary: taken with camera above and below eye level, with one hand or two hands. Pitch angles vary more than yaw angles, which are frontal. Some perspective distortion is evident.
- ▷ The images have mean IOD of 140 pixels.
- ▷ The images are of subjects from the United States.
- ▷ The images are of adults.
- ▷ The images are all live capture.
- ▷ When these images are input to the algorithm, they are labelled as being of type ”WILD” - see Table 4 of the FRVT API.

4.5 Webcam images

- ▷ The number of images is below 1500.
- ▷ The number of subjects is below 1500.
- ▷ All subjects have a webcam image, and a portrait image.
- ▷ The portrait images are in reasonable conformance with the ISO/IEC 19794-5 Full Frontal image type.
- ▷ The webcam images are taken with camera at a typical head height, with mild pitch angles, low yaw angles, but some variation in range, such that low perspective distortion is sometimes evident.
- ▷ The images have mean IOD of 68 pixels (sd=12).
- ▷ The images are of subjects from the United States.
- ▷ The images are of adults.
- ▷ The images are all live capture.
- ▷ When these images are input to the algorithm, they are labelled as being of type ”MUGSHOT” - see Table 4 of the FRVT API.

(a) Visa (b) Mugshot (c) Webcam (d) Selfie (e) Wild

Figure 3: The figure gives simulated samples of image types used in this report.

4.6 Wild images

- ▷ The number of images is O(10^5).
- ▷ The number of subjects is O(10^3).
- ▷ The number of subjects with two images O(10^3).
- ▷ The images include many photojournalism-style images. Images are given to the algorithm using a variable but generally tight crop of the head. Resolution varies very widely. The images are very unconstrained, with wide yaw and pitch pose variation. Faces can be occluded, including hair and hands.
- ▷ The images are of adults.
- ▷ All of the images are live capture, none are scanned.
- ▷ When these images are input to the algorithm, they are labelled as being of type "WILD" - see Table 4 of the FRVT API.

5 Results

5.1 Test goals

- ▷ To state overall accuracy.
- ▷ To compare algorithms.

5.2 Test design

Method: For visa images:

- ▷ The comparisons are of visa photos against visa photos.
- ▷ The number of genuine comparisons is O(10^4).
- ▷ The number of impostor comparisons is O(10^{10}).

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

- ▷ The comparisons are fully zero-effort, meaning impostors are paired without attention to sex, age or other covariates. However, later analysis is conducted on subsets.
- ▷ The number of persons is O(10^5).
- ▷ The number of images used to make 1 template is 1.
- ▷ The number of templates used to make each comparison score is two corresponding to simple one-to-one verification.

For mugshot images:

- ▷ The comparisons are of mugshot photos against mugshot photos.
- ▷ The number of genuine comparisons is O(10^5).
- ▷ The number of impostor comparisons is O(10^7).
- ▷ The comparisons are fully zero-effort, meaning impostors are paired without attention to sex, age or other covariates.
- ▷ The number of persons is O(10^6).
- ▷ The number of images used to make 1 template is 1.
- ▷ The number of templates used to make each comparison score is two corresponding to simple one-to-one verification.

For selfie images:

- ▷ The comparisons are of selfie photos against portrait photos.
- ▷ The number of genuine comparisons is O(10^2).
- ▷ The number of impostor comparisons is O(10^8) selfies are compared with portraits of O(10^6) other subjects.
- ▷ The comparisons are fully zero-effort, meaning impostors are paired without attention to sex, age or other covariates.
- ▷ The number of persons is O(10^6).
- ▷ The number of images used to make 1 template is 1.
- ▷ The number of templates used to make each comparison score is two corresponding to simple one-to-one verification.

For webcam images:

- ▷ The comparisons are of webcam photos against portrait photos.
- ▷ The number of genuine comparisons is O(10^3).
- ▷ The number of impostor comparisons is O(10^9) webcams are compared with portraits of O(10^6) other subjects.
- ▷ The comparisons are fully zero-effort, meaning impostors are paired without attention to sex, age or other covariates.
- ▷ The number of persons is O(10^6).
- ▷ The number of images used to make 1 template is 1.

- ▷ The number of templates used to make each comparison score is two corresponding to simple one-to-one verification.

For child exploitation images:

- ▷ The comparisons are of unconstrained child exploitation photos against others of the same type.
- ▷ The number of genuine comparisons is O(10^4).
- ▷ The number of impostor comparisons is O(10^7).
- ▷ The comparisons are fully zero-effort, meaning impostors are paired without attention to sex, age or other covariates.
- ▷ The number of persons is O(10^3).
- ▷ The number of images used to make 1 template is 1.
- ▷ The number of templates used to make each comparison score is two corresponding to simple one-to-one verification.
- ▷ We produce two performance statements. First, is a DET as used for visa and mugshot images. The second is a cumulative match characteristic (CMC) summarizing a simulated one-to-many search process. This is done as follows.
  - We regard M enrollment templates as items in a gallery.
  - These M templates come from $M > N$ individuals, because multiple images of a subject are present in the gallery under separate identifiers.
  - We regard the verification templates as search templates.
  - For each search we compute the rank of the highest scoring mate.
  - This process should properly be conducted with a 1:N algorithm, such as those tested in NIST IR 8009. We use the 1:1 algorithms in a simulated 1:N mode here to a) better reflect what a child exploitation analyst does, and b) to do show algorithm efficacy is better than that revealed in the verification DETs.

5.3 Failure to enrol

Algorithm	Failure to Enrol Rate[1]											
Name	CHILD-EXPLOIT		MUGSHOT		SELFIES		VISA		WEBCAM		WILD	
3divi-000	0.2019	18	0.0019	22	0.0202	18	0.0008	18	0.0020	15	0.2070	25
3divi-001	0.2006	17	0.0019	21	-	39	0.0007	17	-	39	0.1898	23
ayonix-000	0.0000	1	0.0109	37	0.0751	30	0.0137	42	0.0109	22	0.0000	2
camvi-001	0.3931	32	0.0033	27	0.0231	19	0.0010	20	0.0027	17	0.3767	32
cyberextruder-001	0.5338	38	0.0036	28	0.0376	21	0.0029	31	0.0205	27	0.5730	39
dermalog-002	0.9109	44	0.0045	32	0.0954	33	0.0013	22	0.0471	31	0.3979	33
dermalog-003	0.0434	9	0.0007	9	0.0000	1	0.0025	29	0.0007	6	0.0701	14
digitalbarriers-000	0.5469	39	0.0043	30	0.0925	31	0.0019	27	0.0184	26	0.5170	38
digitalbarriers-001	0.5102	37	0.0044	31	0.0925	32	0.0018	26	0.0232	28	0.4593	37
id3-001	0.3411	27	0.0043	29	0.0260	20	0.0043	39	0.0014	12	0.2090	26
id3-002	0.3168	24	0.0030	26	0.0202	17	0.0032	33	0.0020	14	0.1496	16
innovatrics-000	0.3392	25	0.0013	14	0.0087	12	0.0004	10	0.0027	16	0.3337	30
innovatrics-001	0.3392	26	0.0013	15	0.0087	15	0.0004	11	0.0027	18	0.3337	31
isityou-000	0.4714	35	0.0022	25	0.0665	28	0.0010	19	0.0116	24	0.4586	36
itmo-001	0.5751	42	0.0114	38	0.0694	29	0.0047	40	0.0546	33	0.7410	43
itmo-002	0.5751	43	0.0068	35	0.0636	27	0.0029	32	0.0498	32	0.6985	41
morpho-000	0.0000	4	0.0000	2	0.0000	4	0.0000	3	0.0000	4	0.0000	6
morpho-001	0.0000	6	0.0000	3	-	39	0.0000	5	-	39	0.0000	7
neurotechnology-000	0.0000	5	0.0000	5	0.0000	5	0.0000	4	0.0000	5	0.0163	10
neurotechnology-001	0.2962	22	0.0000	4	0.0000	2	0.0000	1	0.0000	1	0.0163	9
noblis-000	0.0000	2	0.0000	1	-	39	0.0000	2	-	39	0.0000	3
ntechlab-001	0.0926	15	0.0009	11	0.0029	7	0.0005	13	0.0007	8	0.0584	13
ntechlab-002	0.0926	14	0.0009	10	-	39	0.0005	12	-	39	0.0584	12
rankone-000	0.0187	8	0.0005	7	0.0000	6	0.0003	8	0.0007	10	0.2349	27
rankone-002	0.0009	7	0.0001	6	0.0000	3	0.0000	6	0.0000	3	0.0855	15
samtech-000	0.5474	40	0.0052	33	0.0491	25	0.0042	38	0.0252	29	0.7023	42
tongyitrans-001	0.0000	3	0.0068	34	0.0462	23	0.0040	36	0.0055	20	0.0000	4
tongyitrans-002	0.3609	30	0.0078	36	0.0462	24	0.0040	37	0.0055	21	0.0000	5
vcog-001	0.1579	16	-	39	0.0058	10	0.0018	25	0.0000	2	-	39
vcog-002	0.2209	19	0.0021	23	0.0087	14	0.0019	28	0.0007	9	0.1672	18
vigilantsolutions-000	0.5580	41	0.0018	20	0.0462	22	0.0007	16	0.0109	23	0.5927	40
vigilantsolutions-001	0.3585	29	0.0010	13	0.0116	16	0.0004	9	0.0048	19	0.3262	29
visionlabs-001	0.2699	21	0.0014	16	0.0058	11	0.0014	24	0.0020	13	0.1803	21
visionlabs-002	0.3063	23	0.0021	24	0.0087	13	0.0026	30	0.0007	7	0.1865	22
vocord-001	0.4732	36	0.0158	39	0.0520	26	0.0038	35	0.0348	30	0.4494	35
vocord-002	0.3782	31	0.0015	17	0.0029	9	0.0037	34	0.0171	25	0.1992	24
yisheng-000	0.4277	34	0.0016	19	-	39	0.0006	15	-	39	0.1770	20
yisheng-001	0.4277	33	-	39	-	39	0.0006	14	-	39	0.1770	19
yitu-000	0.3475	28	0.0015	18	0.0029	8	0.0013	23	0.0014	11	0.1591	17

Table 3: FTE is the proportion of failed template generation attempts. Failures can occur because the software throws an exception, or because the software electively refuses to process the input image. This would typically occur if a face is not detected. FTE is measured as the number of function calls that give a non-zero error code, OR that give a "small" template. This is defined as one whose size is less than 0.3 times the median template size. This second rule is needed because some algorithms incorrectly fail to return a non-zero error code when template generation fails.

[1] *The effects of FTE are included in the accuracy results of this report by regarding any template comparison involving a failed template enrollment to produce a low similarity score. Thus higher FTE results in higher FNMR.*

5.4 Recognition accuracy

Core algorithm accuracy is stated via:

- ▷ **Cooperative subjects**
  - The summary table of Figure 2;
  - The visa image DETs of Figures 4 and 5;
  - The mugshot DETs of Figure 6 ;
  - The selfie-portrait DETs of Figure 7;
  - The webcam-portrait DETs of Figure 8;
- ▷ **Cooperative subjects**
  - The photojournalism DET of Figure 9
  - The child-exploitation DET of Figure 10;
  - The child-exploitation CMC of Figure 11.

Figure 15 shows dependence of false match rate on algorithm score threshold. This allows a deployer to set a threshold to target a particular false match rate appropriate to the security objectives of the application.

Figure 17 likewise shows FMR(T) but for mugshots, and specially four subsets of the population.

Note that in both the mugshot and visa sets false match rates vary with the ethnicity, age, and sex, of the enrollee and impostor - see section 5.6.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 4: For the visa images, detection error tradeoff (DET) characteristics showing false non-match rate vs. false match rate plotted parametrically on threshold, T. The scales are logarithmic in order to show many decades of FMR.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 5: For the visa images, but now just for low FMR*, detection error tradeoff* (DET) *characteristics showing false non-match rate vs. false match rate plotted parametrically on threshold, T. The scales are logarithmic in order to show many decades of* FMR.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 6: For the mugshot images, detection error tradeoff (DET) characteristics showing false non-match rate vs. false match rate plotted parametrically on threshold, T. The scales are logarithmic in order to show decades of FMR.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 7: For the selfie-to-portrait comparisons, detection error tradeoff (DET) characteristics showing false non-match rate vs. false match rate plotted parametrically on threshold, T. The scales are logarithmic in order to show several decades of FMR. Caution: The FNMR values here are optimistic statements of accuracy because the image pairs were collected on the same day. This is known across biometrics to give better accuracy, and is operationally relevant only in special cases.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 8: For the webcam-to-portrait comparisons, detection error tradeoff (DET) characteristics showing false non-match rate vs. false match rate plotted parametrically on threshold, T. The scales are logarithmic in order to show several decades of FMR. Caution: The FNMR values here are optimistic statements of accuracy because the image pairs were collected on the same day. This is known across biometrics to give better accuracy, and is operationally relevant only in special cases.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 9: For the wild image comparisons, detection error tradeoff (DET) characteristics showing false non-match rate vs. false match rate plotted parametrically on threshold, T. The scales are logarithmic in order to show several decades of FMR.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 10: For child exploitation images, detection error tradeoff (DET) characteristics showing false non-match rate vs. false match rate plotted parametrically on threshold, T. The scales are logarithmic in order to show many decades of FMR. Accuracy is poor because many images have adverse quality characteristics, and because detection and enrollment fails.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 11: For child exploitation images, cumulative match characteristics (CMC) showing true positive identification rate vs. rank. This is simulation of a one-to-many search experiment - see discussion in section 5.2. The scales are logarithmic in order to show the effect of long candidate lists. Accuracy is poor but much improved relative to the 1:1 DETs of Fig. 10 because a search can succeed if any of a subject's several enrolled images matches the search image with a high score.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 12: For the mugshot images, error tradeoff characteristics for white females, black females, black males and white males. The grey lines correspond to fixed thresholds, showing how both FNMR and FMR vary at one operating threshold. Important: Many of the plots will naively be read as saying whites gives lower error rates than blacks because the blue traces lie beneath the red ones. However, this is misleading and incomplete: The grey lines show the traces are generally shifted horizontally. Thus for the dermalog-001 algorithm FNMR for whites is higher than for blacks at a fixed threshold but, at the same time, FMR is higher for blacks - see Figure 17. As access control systems almost always operate at a fixed threshold, the naive interpretation is incorrect.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 13: For the visa images, FNMR and FMR at six operating points along the DET characteristic. At each point a line is drawn between (FMR,FNMR)$_{\text{MALE}}$ and (FMR,FNMR)$_{\text{FEMALE}}$ showing how which sex has lower FMR and/or FNMR. The "M" label denotes male, the other end of the line corresponds to female. The six operating thresholds are selected to give the nominal false match rates given in the legend, and are computed over all impostor pairs regardless of age, sex, and place of birth. The plotted FMR values are broadly an order of magnitude larger than the nominal rates because FMR is computed over demographically-matched impostor pairs i.e individuals of the same sex, from the same geographic region (see section 5.6.1), and the same age group (see section 5.6.2).

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 14: For the wild image comparisons, detection error tradeoff (DET) characteristics showing false non-match rate vs. false match rate plotted parametrically on threshold, T. Error rates are higher here than in the generic wild DET (Fig 9) becuase the impostor pairs here are same-sex only. The scales are logarithmic in order to show several decades of FMR.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 15: For the visa images, the false match calibration curves show false match rate vs. threshold. These curves apply to zero-effort impostors. As shown later (sec. 5.6), FMR is higher for demographic-matched impostors.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 16: For the visa images, the curves show how false matches are concentrated in certain images. Specifically each line plots $FMR(k)$ with k the number of images rejected in decreasing order of how many false matches that image was involved in. $FMR(0) = 10^{-4}$. In terms of the biometric zoo, the most "wolf-ish" images are rejected first i.e. those enrollment or verification images most often involved in false matches. A flatter response is considered superior. A steeply descending response indicates that certain kinds of images false match against others, e.g. if hypothetically images of men with particular mustaches would falsely match others.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 17: For the mugshot images, the false match calibration curves show false match rate vs. threshold. Separate curves appear for white females, black females, black males and white males.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

5.5 Genuine distribution stability

5.5.1 Effect of birth place on the genuine distribution

Background: Both skin tone and bone structure vary geographically. Prior studies have reported variations in FNMR and FMR.

Goal: To measure false non-match rate (FNMR) variation with country of birth.

Methods: Thresholds are determined that give FMR = $\{0.001, 0.0001\}$ over the entire impostor set. Then FNMR is measured over 1000 bootstrap replications of the genuine scores. Only those countries with at least 140 individuals are included in the analysis.

Results: Figure 18 shows FNMR by country of birth for the two thresholds.

Caveats: The results may not relate to subject-specific properties. Instead they could reflect image-specific quality differences, which could occur due to collection protocol or software processing variations.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 18: For the visa images, the dots show FNMR by country of birth for two operating thresholds corresponding to FMR = $\{0.001, 0.0001\}$ computed over all $O(10^{10})$ impostor scores. The figures shows an order of magnitude variation in FNMR across country of birth; these effects are due to quality variations. The least accurate countries vary by algorithm.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

5.5.2 Effect of age on genuine subjects

Background: Faces change appearance throughout life. Face recognition algorithms have previously beeen reported to give better accuracy on older individuals (See NIST IR 8009).

Goal: To quantify false non-match rates (FNMR) as a function of age. We do not aim to quantify ageing effects here as the separation between two samples is limited to just a few years.

Methods: Using the visa images, thresholds are determined that give FMR = 0.001 and 0.0001 over the entire impostor set. Then FNMR is measured over 1000 bootstrap replications of the genuine scores.

Results: For the visa images, Figure 19 shows how false non-match rates for genuine users, as a function of age group.

The notable aspects are:

- ▷ Younger subjects give considerably higher FNMR. This is likely due to rapid growth and change in facial appearance.
- ▷ FNMR trends down throughout life. The last bin, AGE $>$ 72, contains fewer than 140 mated pairs, and may be affected by small sample size.

Caveats: None.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 19: For the visa images, the dots show FNMR by age group for two operating thresholds corresponding to FMR = $\{0.001, 0.0001\}$ computed over all $O(10^{10})$ impostor scores. Given a pair of face images taken at different times, we assign a false non-match to the bin that is the arithmetic average of the subject's ages. This plot shows only the effect of age, not ageing. The number of comparisons in each bin is generally in the thousands. However the FNMR for the first and last bins are each computed over fewer than 150 comparisons.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

5.6 Impostor distribution stability

5.6.1 Effect of birth place on the impostor distribution

Background: Facial appearance varies geographically, both in terms of skin tone, cranio-facial structure and size. This section addresses whether false match rates vary intra- and inter-regionally.

Goals:

- ▷ To show the effect of birth region of the impostor and enrollee on false match rates.
- ▷ To determine whether some algorithms give better impostor distribution stability.

Methods:

- ▷ For the visa images, NIST defined 10 regions: Sub-Saharan Africa, South Asia, Polynesia, North Africa, Middle East, Europe, East Asia, Central and South America, Central Asia, and the Caribbean.
- ▷ For the visa images, NIST mapped each country of birth to a region. There is some arbitrariness to this. For example, Egypt could reasonably be assigned to the Middle East instead of North Africa. An alternative methodology could, for example, assign the Philippines to *both* Polynesia and East Asia.
- ▷ FMR is computed for cases where all face images of impostors born in region r_2 are compared with enrolled face images of persons born in region r_1.

$$\text{FMR}(r_1, r_2, T) = \frac{\sum_{i=1}^{N_{r_1,r_2}} H(s_i - T)}{N_{r_1,r_2}} \tag{5}$$

  where the same threshold, T, is used in all cells, and H is the unit step function. The threshold is set to give FMR(T) = 0.001 over the entire set of visa image impostor comparisons.
- ▷ This analysis is then repeated by country-pair, but only for those country pairs where both have at least 1000 images available. The countries[1] appear in the axes of graphs that follow.
- ▷ The mean number of impostor scores in any cross-region bin is 33 million. The smallest number of impostor scores in any bin is 135000, for Central Asia - North Africa. While these counts are large enough to support reasonable significance, the number of individual faces is much smaller, O($N^{0.5}$).
- ▷ The numbers of impostor scores in any cross-country bin is shown in Figure 87.

Results: Subsequent figures show heatmaps that use color to represent the base-10 logarithm of the false match rate. Red colors indicate high (bad) false match rates. Dark colors indicate benign false match rates. There are two series of graphs corresponding to aggregated geographical regions, and to countries. The notable observations are:

- ▷ The on-diagonal elements correspond to within-region impostors. FMR is generally above the nominal value of FMR = 0.001. Particularly there is usually higher FMR in, Sub-Saharan Africa, South Asia, and the Caribbean. Europe and Central Asia, on the other hand, usually give FMR closer to the nominal value.
- ▷ The off-diagonal elements correspond to across-region impostors. The highest FMR is produced between the Caribbean and Sub-Saharan Africa.
- ▷ Algorithms vary.

[1] These are Argentina, Australia, Brazil, Chile, China, Costa Rica, Cuba, Czech Republic, Dominican Republic, Ecuador, Egypt, El Salvador, Germany, Ghana, Great Britain, Greece, Guatamala, Haiti, Hong Kong, Honduras, Indonesia, India, Israel, Jamaica, Japan, Kenya, Korea, Lebanon, Mexico, Malaysia, Nepal, Nigeria, Peru, Philippines, Pakistan, Poland, Romania, Russia, South Africa, Saudi Arabia, Thailand, Trinidad, Turkey, Taiwan, Ukraine, Venezuela, and Vietnam.

- ▷ We computed the same quantities for a global FMR = 0.0001. The effects are similar.

Caveats:

- ▷ The effects of variable impostor rates on one-to-many identification systems may well differ from what's implied by these one-to-one verification results. Two reasons for this are a) the enrollment galleries are usually imbalanced across countries of birth, age and sex; b) one-to-many identification algorithms often implement techniques aimed at stabilizing the impostor distribution. Further research is necessary.
- ▷ In principle, the effects seen in this subsection could be due to differences in the image capture process. We consider this unlikely since the effects are maintained across geography - e.g. Caribbean vs. Africa, or Japan vs. China.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 20: For visa images, the heatmap shows how the mean of the impostor distribution for the country pair (a,b) is shifted relative to the mean of the global impostor distribution, expressed as a number of standard deviations of the global impostor distribution. This statistic is designed to show shifts in the entire impostor distribution, not just tail effects that manifest as the anomalously high (or low) false match rates that appear in the subsequent figures. The countries are chosen to show that skin tone alone does not explain impostor distribution shifts. The reduced shift in Asian populations with the Yitu and TongYiTrans algorithms, is accompanied by positive shifts in the European populations. This reversal relative to most other algorithms, may derive from use of nationally weighted training sets. The Visionlabs algorithm appears most insensitive to country effects. The figure is computed from same-sex and same-age impostor pairs.

Figure 21: For algorithm 3divi-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 22: For algorithm 3divi-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 23: For algorithm camvi-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 24: For algorithm cyberextruder-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 25: For algorithm dermalog-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 26: For algorithm dermalog-003 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 27: For algorithm digitalbarriers-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 28: For algorithm digitalbarriers-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 29: For algorithm id3-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 30: For algorithm id3-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 31: For algorithm isityou-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 32: For algorithm morpho-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 33: For algorithm morpho-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 34: For algorithm neurotechnology-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 35: For algorithm neurotechnology-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 36: For algorithm noblis-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 37: For algorithm ntechlab-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 38: For algorithm ntechlab-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 39: For algorithm rankone-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 40: For algorithm rankone-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 41: For algorithm samtech-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 42: For algorithm tongyitrans-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 43: For algorithm tongyitrans-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 44: For algorithm vcog-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 45: For algorithm vcog-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 46: For algorithm vigilantsolutions-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 47: For algorithm vigilantsolutions-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 48: For algorithm visionlabs-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 49: For algorithm visionlabs-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 50: For algorithm vocord-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 51: For algorithm yisheng-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 52: For algorithm yisheng-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 53: For algorithm yitu-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given region pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 54: For algorithm 3divi-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 55: For algorithm 3divi-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 56: For algorithm camvi-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 57: For algorithm cyberextruder-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 58: For algorithm dermalog-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 59: For algorithm dermalog-003 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 60: For algorithm digitalbarriers-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 61: For algorithm digitalbarriers-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 62: For algorithm id3-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 63: For algorithm id3-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 64: For algorithm isityou-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 65: For algorithm morpho-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 66: For algorithm morpho-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 67: For algorithm neurotechnology-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 68: For algorithm neurotechnology-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 69: For algorithm noblis-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 70: For algorithm ntechlab-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 71: For algorithm ntechlab-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 72: For algorithm rankone-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 73: For algorithm rankone-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 74: For algorithm samtech-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 75: For algorithm tongyitrans-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 76: For algorithm tongyitrans-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 77: For algorithm vcog-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 78: For algorithm vcog-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 79: For algorithm vigilantsolutions-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 80: For algorithm vigilantsolutions-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 81: For algorithm visionlabs-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 82: For algorithm visionlabs-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 83: For algorithm vocord-002 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 84: For algorithm yisheng-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 85: For algorithm yisheng-001 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 86: For algorithm yitu-000 operating on visa images, the heatmap shows false match rates observed over impostor comparisons of faces from different individuals who were born in the given country pair. False matches are counted against a recognition threshold fixed globally to give the target FMR in the plot title, computed over all $O(10^{10})$ impostor comparisons. If text appears in each box it give the same quantity as that coded by the color. Grey indicates FMR is at the intended FMR target level. Light red colors present a security vulnerability to, for example, a passport gate. Each +1 increase in log10 FMR corresponds to a factor of 10 increase in FMR. The matrix is not quite symmetric because images in the enrollment and verification sets are different.

Figure 87: For visa images, the heatmap shows the count of impostor comparisons of faces from different individuals who were born in the given country pair.

5.6.2 Effect of age on impostors

Background: This section shows the effect of age on the impostor distribution. The ideal behaviour is that the age of the enrollee and the impostor would not affect impostor scores. This would support FMR stability over sub-populations.

Goals:

- ▷ To show the effect of relative ages of the impostor and enrollee on false match rates.
- ▷ To determine whether some algorithms have better impostor distribution stability.

Methods:

- ▷ Define 14 age group bins, spanning 0 to over 100 years old.
- ▷ Compute FMR over all impostor comparisons for which the subjects in the enrollee and impostor images have ages in two bins.
- ▷ Compute FMR over all impostor comparisons for which the subjects are additionally of the same sex, and born in the same geographic region.

Results:

The notable aspects are:

- ▷ Diagonal dominance: Impostors are more likely to be matched against their same age group.
- ▷ Same sex and same region impostors are more successful. On the diagonal, an impostor is more likely to succeed by posing as someone of the same sex. If $\Delta \log_{10} \text{FMR} = 0.2$, then same-sex same-region FMR exceeds the all-pairs FMR by factor of $10^{0.2} = 1.6$.
- ▷ Young children impostors give elevated FMR against young children. Older adult impostor give elevated FMR against older adults. These effects are quite large, for example if $\Delta \log_{10} \text{FMR} = 1.0$ larger than a 32 year old, then these groups have higher FMR by a factor of $10^1 = 10$. This would imply an FMR above 0.01 for a nominal (global) FMR = 0.001.
- ▷ Algorithms vary.
- ▷ We computed the same quantities for a global FMR = 0.0001. The effects are similar.

Note the calculations in this section include impostors paired across all countries of birth.

Figure 88: For algorithm 3divi-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 89: For algorithm 3divi-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 90: For algorithm camvi-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 91: For algorithm cyberextruder-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 92: For algorithm dermalog-002 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 93: For algorithm dermalog-003 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 94: For algorithm digitalbarriers-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 95: For algorithm digitalbarriers-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 96: For algorithm id3-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 97: For algorithm id3-002 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 98: For algorithm isityou-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 99: For algorithm morpho-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 100: For algorithm morpho-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 101: For algorithm neurotechnology-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 102: For algorithm neurotechnology-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 103: For algorithm noblis-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 104: For algorithm ntechlab-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 105: For algorithm ntechlab-002 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 106: For algorithm rankone-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 107: For algorithm rankone-002 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 108: For algorithm samtech-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 109: For algorithm tongyitrans-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 110: For algorithm tongyitrans-002 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 111: For algorithm vcog-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 112: For algorithm vcog-002 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 113: For algorithm vigilantsolutions-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 114: For algorithm vigilantsolutions-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 115: For algorithm visionlabs-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 116: For algorithm visionlabs-002 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 117: For algorithm vocord-002 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 118: For algorithm yisheng-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 119: For algorithm yisheng-001 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Figure 120: For algorithm yitu-000 operating on visa images, the heatmap shows false match observed over impostor comparisons of faces from different individuals who have the given age pair. False matches are counted against a recognition threshold fixed globally to give FMR = 0.001 over all $O(10^{10})$ impostor comparisons. The text in each box gives the same quantity as that coded by the color. Light colors present a security vulnerability to, for example, a passport gate.

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

Accuracy Terms + Definitions

In biometrics, Type II errors occur when two samples of one person do not match – this is called a **false negative**. Correspondingly, Type I errors occur when samples from two persons do match – this is called a **false positive**. Matches are declared by a biometric system when the native comparison score from the recognition algorithm meets some **threshold**. Comparison scores can be either **similarity scores**, in which case higher values indicate that the samples are more likely to come from the same person, or **dissimilarity scores**, in which case higher values indicate different people. Similarity scores are traditionally computed by **fingerprint** and **face** recognition algorithms, while dissimilarities are used in **iris recognition**. In some cases, the dissimilarity score is a distance; this applies only when **metric** properties are obeyed. In any case, scores can be either **mate** scores, coming from a comparison of one person's samples, or **nonmate** scores, coming from comparison of different persons' samples. The words **genuine** or **authentic** are synonyms for mate, and the word **impostor** is used a synonym for nonmate. The words mate and nonmate are traditionally used in identification applications (such as law enforcement search, or background checks) while genuine and impostor are used in verification applications (such as access control).

A **error tradeoff** characteristic represents the tradeoff between Type II and Type I classification errors. For verification this plots false non-match rate (FNMR) vs. false match rate (FMR) parametrically with T.

The error tradeoff plots are often called **detection error tradeoff (DET)** characteristics or **receiver operating characteristic** (ROC). These serve the same function but differ, for example, in plotting the complement of an error rate (e.g. TMR = 1 – FNMR) and in transforming the axes most commonly using logarithms, to show multiple decades of FMR. More rarely, the function might be the inverse Gaussian function.

More detail and generality is provided in formal biometrics testing standards, see the various parts of ISO/IEC 19795 Biometrics Testing and Reporting. More terms, including and beyond those to do with accuracy, see ISO/IEC 2382-37 Information technology -- Vocabulary -- Part 37: Harmonized biometric vocabulary

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

FNMR(T) "False non-match rate"
FMR(T) "False match rate"

FNMR(T) "False non-match rate"
FMR(T) "False match rate"